

cartpole_swingup.py

```

1  """
2  Starter code for the problem "Cart-pole swing-up".
3
4  Autonomous Systems Lab (ASL), Stanford University
5  """
6
7  import time
8
9  from animations import animate_cartpole
10
11 import jax
12 import jax.numpy as jnp
13
14 import matplotlib.pyplot as plt
15
16 import numpy as np
17
18 from scipy.integrate import odeint
19
20
21 def linearize(f, s, u):
22     """Linearize the function `f(s, u)` around `(s, u)`.
23
24     Arguments
25     -----
26     f : callable
27         A nonlinear function with call signature `f(s, u)`.
28     s : numpy.ndarray
29         The state (1-D).
30     u : numpy.ndarray
31         The control input (1-D).
32
33     Returns
34     -----
35     A : numpy.ndarray
36         The Jacobian of `f` at `(s, u)`, with respect to `s`.
37     B : numpy.ndarray
38         The Jacobian of `f` at `(s, u)`, with respect to `u`.
39     """
40     # WRITE YOUR CODE BELOW #####
41     # INSTRUCTIONS: Use JAX to compute `A` and `B` in one line.
42     A,B = jax.jacobian(f, argnums=(0,1))(s,u)
43     #####
44     return A, B
45
46
47 def ilqr(f, s0, s_goal, N, Q, R, QN, eps=1e-3, max_iters=1000):
48     """Compute the iLQR set-point tracking solution.

```

```

49
50 Arguments
51 -----
52 f : callable
53     A function describing the discrete-time dynamics, such that
54     `s[k+1] = f(s[k], u[k])`.
55 s0 : numpy.ndarray
56     The initial state (1-D).
57 s_goal : numpy.ndarray
58     The goal state (1-D).
59 N : int
60     The time horizon of the LQR cost function.
61 Q : numpy.ndarray
62     The state cost matrix (2-D).
63 R : numpy.ndarray
64     The control cost matrix (2-D).
65 QN : numpy.ndarray
66     The terminal state cost matrix (2-D).
67 eps : float, optional
68     Termination threshold for iLQR.
69 max_iters : int, optional
70     Maximum number of iLQR iterations.
71
72 Returns
73 -----
74 s_bar : numpy.ndarray
75     A 2-D array where `s_bar[k]` is the nominal state at time step `k`,
76     for `k = 0, 1, ..., N-1`
77 u_bar : numpy.ndarray
78     A 2-D array where `u_bar[k]` is the nominal control at time step `k`,
79     for `k = 0, 1, ..., N-1`
80 Y : numpy.ndarray
81     A 3-D array where `Y[k]` is the matrix gain term of the iLQR control
82     law at time step `k`, for `k = 0, 1, ..., N-1`
83 y : numpy.ndarray
84     A 2-D array where `y[k]` is the offset term of the iLQR control law
85     at time step `k`, for `k = 0, 1, ..., N-1`
86 """
87 if max_iters <= 1:
88     raise ValueError("Argument `max_iters` must be at least 1.")
89 n = Q.shape[0] # state dimension
90 m = R.shape[0] # control dimension
91
92 # Initialize gains `Y` and offsets `y` for the policy
93 Y = np.zeros((N, m, n))
94 y = np.zeros((N, m))
95
96 # Initialize the nominal trajectory `(s_bar, u_bar)`, and the
97 # deviations `(ds, du)`
98 u_bar = np.zeros((N, m))

```

```

99     s_bar = np.zeros((N + 1, n))
100    s_bar[0] = s0
101    for k in range(N):
102        s_bar[k + 1] = f(s_bar[k], u_bar[k])
103    ds = np.zeros((N + 1, n))
104    du = np.zeros((N, m))
105
106    # iLQR loop
107    converged = False
108    for _ in range(max_iters):
109        # Linearize the dynamics at each step `k` of `(s_bar, u_bar)`
110        A, B = jax.vmap(linearize, in_axes=(None, 0, 0))(f, s_bar[:-1], u_bar)
111        A, B = np.array(A), np.array(B)
112
113        # PART (c) #####
114        # INSTRUCTIONS: Update `Y`, `y`, `ds`, `du`, `s_bar`, and `u_bar`.
115        Vxx = QN
116        Vx = QN@(s_bar[N, :]-s_goal)
117        V = (s_bar[N, :]-s_goal).T@QN@(s_bar[N, :]-s_goal)
118        for k in reversed(range(N)):
119            Qk = (s_bar[k, :]-s_goal).T@Q@(s_bar[k, :]-s_goal) + (u_bar[k, :]).T@R@(u_bar[k, :])
120        + V
121            Qx = Q@(s_bar[k, :]-s_goal) + A[k, :, :].T@Vx
122            Qu = R@(u_bar[k, :]) + B[k, :, :].T@Vx
123            Qxx = Q + A[k, :, :].T@Vxx@A[k, :, :]
124            Quu = R + B[k, :, :].T@Vxx@B[k, :, :]
125            Qux = B[k, :, :].T@Vxx@A[k, :, :]
126            Y[k, :, :] = -np.linalg.inv(Quu)@Qux
127            y[k, :] = -np.linalg.inv(Quu)@Qu
128            # u_bar[k, :] = u_bar[k, :] + du[k, :]
129            V = Qk - 1/2*y[k, :].T@Quu@y[k, :]
130            Vx = Qx - Y[k, :, :].T@Quu@y[k, :]
131            Vxx = Qxx - Y[k, :, :].T@Quu@Y[k, :, :]
132
133        for k in range(N):
134            du[k, :] = Y[k, :, :](ds[k, :]) + y[k, :]
135            ds[k+1, :] = f(s_bar[k, :], u_bar[k, :]) - s_bar[k+1, :]
136            s_bar[k+1, :] = s_bar[k+1, :] + ds[k+1, :]
137            u_bar[k, :] = u_bar[k, :] + du[k, :]
138
139            #####
140
141        if np.max(np.abs(du)) < eps:
142            converged = True
143            break
144        if not converged:
145            raise RuntimeError("iLQR did not converge!")
146        return s_bar, u_bar, Y, y
147
148    def cartpole(s, u):

```

```

148     """Compute the cart-pole state derivative."""
149     mp = 2.0 # pendulum mass
150     mc = 10.0 # cart mass
151     L = 1.0 # pendulum length
152     g = 9.81 # gravitational acceleration
153
154     x, θ, dx, dθ = s
155     sinθ, cosθ = jnp.sin(θ), jnp.cos(θ)
156     h = mc + mp * (sinθ**2)
157     ds = jnp.array(
158         [
159             dx,
160             dθ,
161             (mp * sinθ * (L * (dθ**2) + g * cosθ) + u[0]) / h,
162             -((mc + mp) * g * sinθ + mp * L * (dθ**2) * sinθ * cosθ + u[0] * cosθ)
163             / (h * L),
164         ]
165     )
166     return ds
167
168
169 # Define constants
170 n = 4 # state dimension
171 m = 1 # control dimension
172 Q = np.diag(np.array([10.0, 10.0, 2.0, 2.0])) # state cost matrix
173 R = 1e-2 * np.eye(m) # control cost matrix
174 QN = 1e2 * np.eye(n) # terminal state cost matrix
175 s0 = np.array([0.0, 0.0, 0.0, 0.0]) # initial state
176 s_goal = np.array([0.0, np.pi, 0.0, 0.0]) # goal state
177 T = 10.0 # simulation time
178 dt = 0.1 # sampling time
179 animate = False # flag for animation
180 closed_loop = True # flag for closed-loop control
181
182 # Initialize continuous-time and discretized dynamics
183 f = jax.jit(cartpole)
184 fd = jax.jit(lambda s, u, dt=dt: s + dt * f(s, u))
185
186 # Compute the iLQR solution with the discretized dynamics
187 print("Computing iLQR solution ... ", end="", flush=True)
188 start = time.time()
189 t = np.arange(0.0, T, dt)
190 N = t.size - 1
191 s_bar, u_bar, Y, y = ilqr(fd, s0, s_goal, N, Q, R, QN)
192 print("done! ({:.2f} s)".format(time.time() - start), flush=True)
193
194 # Simulate on the true continuous-time system
195 print("Simulating ... ", end="", flush=True)
196 start = time.time()
197 s = np.zeros((N + 1, n))

```

```

198 u = np.zeros((N, m))
199 s[0] = s0
200 for k in range(N):
201     # PART (d) #####
202     # INSTRUCTIONS: Compute either the closed-loop or open-loop value of
203     # `u[k]`, depending on the Boolean flag `closed_loop`.
204     if closed_loop:
205         u[k] = Y[k, :, :](s[k, :]-s_bar[k, :])+y[k, :]
206         # raise NotImplementedError()
207     else: # do open-loop control
208         u[k] = u_bar[k]
209         # raise NotImplementedError()
210     #####
211     s[k + 1] = odeint(lambda s, t: f(s, u[k]), s[k], t[k : k + 2])[1]
212 print("done! ({:.2f} s)".format(time.time() - start), flush=True)
213
214 # Plot
215 fig, axes = plt.subplots(1, n + m, dpi=150, figsize=(15, 2))
216 plt.subplots_adjust(wspace=0.45)
217 labels_s = (r"$x(t)$", r"$\theta(t)$", r"$\dot{x}(t)$", r"$\dot{\theta}(t)$")
218 labels_u = (r"$u(t)$",)
219 for i in range(n):
220     axes[i].plot(t, s[:, i])
221     axes[i].set_xlabel(r"$t$")
222     axes[i].set_ylabel(labels_s[i])
223 for i in range(m):
224     axes[n + i].plot(t[:-1], u[:, i])
225     axes[n + i].set_xlabel(r"$t$")
226     axes[n + i].set_ylabel(labels_u[i])
227 if closed_loop:
228     plt.savefig("cartpole_swingup_cl.png", bbox_inches="tight")
229 else:
230     plt.savefig("cartpole_swingup_ol.png", bbox_inches="tight")
231 plt.show()
232
233 if animate:
234     fig, ani = animate_cartpole(t, s[:, 0], s[:, 1])
235     ani.save("cartpole_swingup.mp4", writer="ffmpeg")
236     plt.show()
237

```